

Shriram M

Namakkal, India | shriram.coder@gmail.com | linkedin.com/in/shriram-m-sde | github.com/shriram1206 | [Personal Portfolio](#)

EDUCATION

Selvam College of Technology

Bachelor of Engineering in Computer Science and Engineering

Namakkal, TN

2022 – May 2026

EXPERIENCE

Software Development Engineer Intern

Feb 2026 – Mar 2026

Servimos Technologies Pvt Ltd (WeServeTech LLC)

Chennai, TN

- **Automated weekly compliance reporting** by building an AI report generator in Java 17 and Spring Boot 3.x, cutting manual preparation time from ~4 hours to <10 minutes per cycle.
- **Implemented RAG-based context injection** via pgvector + PostgreSQL for semantic document retrieval and applied MCP (SSE transport, JSON-RPC tool calling) to automate recurring reporting pipelines.
- **Delivered a production-ready AI backend** over 3 Agile sprints with end-to-end ownership from schema design through deployment, with code and architecture reviewed and approved by senior engineers.
- **Established production-grade observability** for the reporting service using structured logging and Spring Boot Actuator metrics, enabling rapid incident debugging and SLA monitoring after go-live.

PROJECTS

Axiom — Multi-Agent RAG Boardroom | *Python, FastAPI, Llama 3.1, ChromaDB, React*

2026 | Source

- **Architected a multi-agent RAG system** that simulates a corporate C-suite, utilizing `asyncio.gather` for concurrent, sub-second orchestration of 3 distinct LLM personas via the Groq API.
- **Engineered strict database-level security isolation** using ChromaDB metadata filtering, ensuring department-scoped knowledge retrieval and zero cross-agent data leakage.
- **Designed a shared-memory boardroom architecture** where agents dynamically evaluate, negotiate, and push back on cross-department proposals based on conflicting system prompts.

Prompt Forge — AI Meta-Framework Synthesis Engine | *Python, FastAPI, Supabase, Groq*

2026 | Live

- **Designed a full-stack SaaS compiler** that transforms vague user inputs into structured, FAANG-grade XML prompts by dynamically evaluating complexity and applying 78+ advanced cognitive frameworks.
- **Integrated Groq's Llama 3 API** via a Python/FastAPI backend, using strict metaprompting and Best-of-N gating to reduce hallucinations while maintaining sub-200ms prompt compilation latency.
- **Hardened API security** by enforcing strict Pydantic payload validation to prevent LLM Denial-of-Wallet attacks, implementing a Debounce API to block disposable emails, and tightening CORS controls.
- **Implemented tier-based usage tracking** on Supabase with RLS and JWT auth to enforce Base vs. ML-Elite rate limits without cross-tenant data leakage.

Quro — AI-Powered Database Assistant | *Java 17, Spring Boot 3.2, PostgreSQL, Docker, Kubernetes*

2026 | Live

- **Built a natural-language-to-SQL engine** using Spring Boot 3.x, the GPT-4o API, PostgreSQL, and Redis caching, enabling non-technical stakeholders to query databases in plain English instead of writing SQL, significantly cutting analytics turnaround time.
- **Improved query accuracy to 92%+** on a test suite of 200+ patterns by iteratively tuning prompts, refining schema context, and adding guardrails around joins, filters, and aggregations.
- **Developed a read-only SQL validation layer** that parses and sanitizes all generated queries, preventing injection attacks, blocking unsafe operations, and enforcing SELECT-only execution on allowed tables.
- **Deployed the engine as a secure REST API** behind Docker/Kubernetes, with observability and audit logging for each natural-language request, generated query, and result set size.

TECHNICAL SKILLS

AI & Agentic Systems: Python, LangChain, Spring AI, RAG Architectures, Agent Orchestration (MCP, asyncio), Prompt Engineering, Llama 3

Backend & APIs: FastAPI, Java 17, Spring Boot 3.x, Node.js, REST/JSON-RPC, SSE Transport

Databases (Vector & SQL): ChromaDB, pgvector, PostgreSQL, Supabase, MySQL

Frontend & Cloud: React.js, Next.js, Tailwind CSS, Docker, Kubernetes, Vercel, Git