

Shriram M

Namakkal, India | shriram.coder@gmail.com | linkedin.com/in/shriram-m-sde | github.com/shriram1206 | [Personal Portfolio](#)

EDUCATION

Selvam College of Technology

Bachelor of Engineering in Computer Science and Engineering

Namakkal, TN

2022 – May 2026

EXPERIENCE

Software Development Engineer Intern

Feb 2026 – Mar 2026

Servimos Technologies Pvt Ltd (WeServeTech LLC)

Chennai, TN

- **Automated weekly compliance reporting** by building an AI report generator in Java 17 and Spring Boot 3.x, cutting manual preparation time from ~4 hours to <10 minutes per cycle.
- **Implemented RAG-based context injection** via pgvector + PostgreSQL for semantic document retrieval and applied MCP (SSE transport, JSON-RPC tool calling) to automate recurring reporting pipelines.
- **Delivered a production-ready AI backend** over 3 Agile sprints with end-to-end ownership from schema design through deployment, with code and architecture reviewed and approved by senior engineers.
- **Established production-grade observability** for the reporting service using structured logging and Spring Boot Actuator metrics, enabling rapid incident debugging and SLA monitoring after go-live.

PROJECTS

Axiom — Multi-Agent RAG Boardroom | Python, FastAPI, Llama 3.1, ChromaDB, React

2026 | Source

- **Architected a multi-agent RAG system** that simulates a corporate C-suite, utilizing `asyncio.gather` for concurrent, sub-second orchestration of 3 distinct LLM personas via the Groq API.
- **Engineered strict database-level security isolation** using ChromaDB metadata filtering, ensuring department-scoped knowledge retrieval and zero cross-agent data leakage.
- **Designed a shared-memory boardroom architecture** where agents dynamically evaluate, negotiate, and push back on cross-department proposals based on conflicting system prompts.

Prompt Forge — AI Meta-Framework Synthesis Engine | Python, FastAPI, Supabase, Groq

2026 | Live

- **Designed a full-stack SaaS compiler** that transforms vague user inputs into structured, FAANG-grade XML prompts by dynamically evaluating complexity and applying 78+ advanced cognitive frameworks.
- **Integrated Groq's Llama 3 API** via a Python/FastAPI backend, using strict metaprompting and Best-of-N gating to reduce hallucinations while maintaining sub-200ms prompt compilation latency.
- **Hardened API security** by enforcing strict Pydantic payload validation to prevent LLM Denial-of-Wallet attacks, implementing a Debounce API to block disposable emails, and tightening CORS controls.
- **Implemented tier-based usage tracking** on Supabase with RLS and JWT auth to enforce Base vs. ML-Elite rate limits without cross-tenant data leakage.

Autonomous Medical Claim Adjudicator Agentic RPA System | UiPath, Gemini 2.5 API, JSON 2026 | Source

- **Integrated Google's Gemini 2.5 Flash LLM** natively into a UiPath architecture via RESTful HTTP protocols to intelligently process unstructured EOB data, bypassing the strict limitations of traditional DOM scraping.
- **Engineered a flawless web traversal loop** using simulated hardware events and fuzzy selectors to autonomously navigate proprietary portals, input patient keys, and securely extract dynamically rendered HTML blocks.
- **Developed advanced string sanitization algorithms** and exact JSON configuration to parse multi-dimensional LLM responses and seamlessly inject AI-adjudicated decision statuses back into local arrays.
- **Hardened the AI automation pipeline** against strict API rate limits and rendering logic using offset delays and `ContinueOnError` suppression to achieve a 100% completion rate without human intervention.

TECHNICAL SKILLS

AI & Agentic Systems: Python, LangChain, Spring AI, RAG Architectures, Agent Orchestration (MCP, asyncio), Prompt Engineering, Llama 3

Backend & APIs: FastAPI, Java 17, Spring Boot 3.x, Node.js, REST/JSON-RPC, SSE Transport

Databases (Vector & SQL): ChromaDB, pgvector, PostgreSQL, Supabase, MySQL

Frontend & Cloud: React.js, Next.js, Tailwind CSS, Docker, Kubernetes, Vercel, Git